Tertiary Infotech Academy Pte Ltd
UEN: 201200696W

ACTIVITY & DEBRIEF PACK

For

WSQ - Generative AI for Problem Solving

TGS Ref No: TGS-2023036653

Conducted by

Tertiary Infotech Academy Pte Ltd
UEN: 201200696W

Version 21.0

TABLE OF CONTENTS

Right-click and choose “Update Field”, or press F9, to build the table of contents.

Activity 1 — Rewriting a Vague Problem into a Measurable Problem Statement

ShopFront SG — 'Customers are unhappy and sales are dropping'

Topic	Topic 1
Objective	A1 · K1 — Identify the type of performance deficiency and express it as a measurable problem statement.
Duration	45 minutes
Grouping	Teams of 3
Tools	ChatGPT / Copilot / Gemini, CollabNote
Ed-tool	CollabNote — https://alfredang.github.io/collabnote/

The Scenario

ShopFront SG is a 42-outlet fashion retailer with an e-commerce site and a mobile app. At the Monday management meeting the Head of Retail Operations opens with one line: "Customers are unhappy and sales are dropping — fix it." No metric, no baseline, no scope. Three departments leave the room with three different interpretations. Marketing books a S$180,000 discount campaign. IT starts an app redesign. Customer Service hires two temps. Ninety days later the complaint rate is unchanged and the budget is gone.

You are the business analyst asked to restate the problem BEFORE any more money is committed. The dashboard extract below is all the evidence you have.

The Evidence

ShopFront SG — 90-day performance extract

Metric	12 months ago	Current	Industry benchmark
Average app page load	2.1 s	4.8 s	under 2.0 s
Checkout completion rate	68%	51%	70%
Cart abandonment	61%	74%	60%
CSAT (post-purchase)	84%	70%	85%
Repeat purchase (90-day)	38%	29%	40%
Complaint tickets / week	120	410	-
In-store sales	S$4.2m/mo	S$4.1m/mo	-
Online sales	S$2.8m/mo	S$1.9m/mo	-

Discussion Questions

1. Which metric in the table is the SYMPTOM the executive noticed, and which metric is closest to the actual performance deficiency? Justify with the numbers.
2. The evidence shows in-store sales are flat while online sales fell 32%. What does that single contrast rule out as a cause, and why does that matter before you spend a dollar?
3. Write the problem statement with all six elements. Which element was hardest to fill, and what does the gap tell you about your organisation's data?
4. Compare your team's statement with the GenAI-generated one. What did the AI assume that is NOT supported by the evidence? (Every team will find at least one.)
5. Whose sign-off do you need on this problem statement before work starts, and what would you show them to get it?

Trainer Debrief

EXPECTED DIRECTION — the deficiency is a DIGITAL CHANNEL conversion failure, not a general 'unhappy customer' problem. In-store being flat while online fell 32% localises the problem to the digital funnel and rules out brand, product and pricing as the primary cause — which immediately invalidates Marketing's S$180k discount campaign.

A defensible statement reads roughly: 'Checkout completion on the ShopFront SG app and web store has fallen from 68% to 51% (benchmark 70%) over 12 months, alongside app load times rising from 2.1s to 4.8s, contributing to a S$0.9m per month decline in online sales. We aim to restore checkout completion to at least 68% and app load to under 2.0s within 90 days, without additional headcount and without changes requiring PDPA re-consent.'

KEY TEACHING POINTS:
1. The executive named a symptom ('unhappy'). CSAT is a lagging perception measure; checkout completion is the operational deficiency you can actually act on.
2. Six elements are non-negotiable. A statement missing a Baseline cannot be evaluated later — this is exactly what Topic 3 comes back to.
3. GenAI reliably invents plausible specifics (a peak-hour window, a payment gateway name, a customer segment) that the evidence never supplied. Learners must mark these as assumptions. This is the single most important AI-literacy habit in the course.
4. Note what has NOT been established yet: WHY load time rose. That is root cause, and it is Activity 2. Resist the team that jumps to 'we need more servers'.

Self-Check

Your statement is complete when a colleague who was NOT in the room can read it and correctly state (a) which metric is being fixed, (b) the number it is at today, (c) the number it must reach, (d) by when, and (e) what you are not allowed to change — without asking you a question. If any of the five is missing, the statement is not finished.

Activity 2 — 5 Whys Root Cause Analysis

Meridian Health Screening — the 07:00 blood-test backlog

Topic	Topic 1
Objective	A2 · K2, K5 — Identify root causes with team members using structured group facilitation.
Duration	45 minutes
Grouping	Teams of 3-5
Tools	5 Whys ed-tool, ChatGPT / Copilot / Gemini
Ed-tool	5 Whys Tool — https://alfredang.github.io/5whys/

The Scenario

Meridian Health operates six health-screening centres in Singapore. Corporate clients book annual screening packages; the fasting blood draw must happen between 07:00 and 09:30.

Over the last four months, 31% of morning appointments overrun by more than 40 minutes. Corporate clients have escalated twice. One major client (1,400 employees, S$310,000 annual contract) has given notice of review. The Operations Manager's proposal is to open a seventh centre at a capital cost of S$1.2m.

The Clinical Director is unconvinced. She has asked your team to establish the ROOT CAUSE before the board approves any capital expenditure. The nurses tell you the same thing every morning: "We're just short-staffed." The data below says something more interesting.

The Evidence

Meridian Health — morning screening operations, 4-month sample

Observation	Value
Booked slots 07:00-09:30	48 per centre
Phlebotomists rostered 07:00	3 of 5
Phlebotomists rostered 08:30	5 of 5
Mean draw time, trained staff	6 min
Mean draw time, relief staff	14 min
Relief staff share of morning roster	40%
Repeat draws (failed first attempt)	18%
Clients arriving without fasting compliance	12%
Lab courier pickup (fixed)	09:45
Samples missing the 09:45 courier	23%
Permanent staff turnover, 12 months	44%

Discussion Questions

6. Run the 5 Whys to five levels. At which level did you stop describing WHAT happens and start explaining WHY the system produces it?
7. The nurses say 'we're short-staffed'. The data shows all 5 phlebotomists are rostered by 08:30. What is the real constraint — headcount, or something else?
8. Trace the causal chain to 44% annual turnover. If turnover is the root cause, what does that say about the S$1.2m seventh centre proposal?
9. Where could this 5 Whys chain have gone wrong? Identify one point where a different, equally plausible 'why' would have led you somewhere completely different.
10. The GenAI facilitator challenged at least one of your answers as opinion rather than fact. Which one, and were you able to supply the evidence?

Trainer Debrief

EXPECTED CHAIN (teams may reach this by slightly different routes — accept any chain that is evidence-supported and reaches a systemic cause):
  Why 1: Appointments overrun -> draws take longer than the slot allows.
  Why 2: Draws take longer -> 40% of morning staff are relief staff at 14 min vs 6 min, and 18% of draws need a repeat.
  Why 3: So many relief staff -> only 3 of 5 permanent phlebotomists are rostered at 07:00.
  Why 4: Under-rostered at 07:00 -> permanent staff turnover is 44%/yr; vacancies are backfilled with relief staff who are not trained to the centre's protocol.
  Why 5: Turnover is 44% -> the early shift is unattractive and there is no structured onboarding or retention path for phlebotomists.

ROOT CAUSE: a workforce retention and onboarding failure, surfacing as a capacity problem.

KEY TEACHING POINTS:
1. The S$1.2m seventh centre solves the SYMPTOM. A new centre staffed by the same 40% untrained relief pool reproduces the same overrun — you buy the problem twice.
2. 'We're short-staffed' was the team's stated cause and it was WRONG in the way that matters: they are not short of headcount, they are short of TRAINED headcount at 07:00. This distinction is worth S$1.2m.
3. Facilitation discipline: one why at a time, and every level must survive 'what evidence supports this?'. The GenAI acting as challenger models the facilitator behaviour learners should copy — this is the K5 group facilitation outcome.
4. Note the branch point: at Why 2 a team could have followed the 12% fasting non-compliance instead and landed on 'client communication'. That is a legitimate SECOND cause. 5 Whys gives depth on one branch — which is exactly why Activity 3 uses Fishbone for breadth.

Self-Check

Your root cause is sound when you can answer YES to all three: (1) If this cause were removed, would the 31% overrun stop recurring — not just improve this month? (2) Is every level in the chain supported by a number in the evidence table, or explicitly flagged as an assumption? (3) Does the root cause point at a SYSTEM (policy, process, incentive) rather than at a person? If your chain ends at 'the relief staff are slow', you have found a symptom and blamed a human — go two levels deeper.

Activity 3 — Fishbone (Ishikawa) Diagram for Breadth of Causes

Horizon Bank — branch service complaints at Jurong East

Topic	Topic 1
Objective	A2 · K2, K5 — Categorise potential causes across all contributing dimensions.
Duration	45 minutes
Grouping	Teams of 3-5
Tools	Fishbone ed-tool, ChatGPT / Copilot / Gemini
Ed-tool	Fishbone Tool — https://alfredang.github.io/fishbone/

The Scenario

Horizon Bank's Jurong East branch has moved from best to worst in the retail network on customer experience in eleven months. Complaints run at 60 per month against a network average of 14. Average counter wait time is 58 minutes at peak; the service standard is 15.

Head Office ran a 'Service Excellence' refresher for all branch staff four months ago. Complaints did not move. The Branch Manager is now under performance review and morale is visibly poor — two tellers have resigned this quarter.

The Regional Director suspects the training was aimed at the wrong cause. Your team has been asked to map ALL contributing causes across every dimension before another intervention is funded. 5 Whys gave depth on one branch; this problem is too broad for one branch.

The Evidence

Horizon Bank Jurong East — observation log and system data

Dimension	Observation
Queue	Single queue for all transaction types; no triage between simple and complex
Peak	11:30-14:00 accounts for 61% of daily footfall (lunch crowd from nearby offices)
Counters	6 counters exist; mean 3.2 staffed at peak
Systems	New KYC platform launched 11 months ago; teller must key client data into 3 systems
Systems	Mean core-banking screen response 8s, was 2s before the KYC platform
Staff	44% of tellers have under 6 months tenure
Staff	Teller KPI is transactions-per-hour; no customer-experience measure
Customer	34% of customers arrive without required documents
Customer	Branch signage and the app's document checklist do not match
Process	Complex cases (loans, disputes) occupy a counter

	for a mean of 27 minutes
Management	No structured feedback route from branch to Head Office product owners
Environment	Only 18 seats in the waiting area; standing customers visibly agitated

Discussion Questions

11. Populate all six bones. Which bone filled up fastest, and which was hardest? What does an empty bone usually mean — no causes, or no visibility?
12. Head Office's answer was a Service Excellence training refresher. Which bone does that intervention target, and what share of your evidenced causes sit on that bone?
13. Find a cause on your diagram that CREATES another cause on a different bone. (Hint: look at the KPI and the queue.) What does that tell you about treating bones as independent?
14. Which three causes would you investigate first? Defend the ranking on evidence strength and likely contribution — not on ease of fixing.
15. The teller KPI is transactions-per-hour. Predict the behaviour that KPI produces at a counter when a confused elderly customer needs 20 minutes. Which bone does that belong on?

Trainer Debrief

EXPECTED DISTRIBUTION — a well-built diagram is heavily loaded on Process, Technology and Measurement, and only lightly on People:

PEOPLE — 44% of tellers under 6 months; two resignations; low morale; thin coaching.

PROCESS — single queue, no triage; complex cases (27 min) block counters; rostering does not match the 11:30-14:00 peak (3.2 of 6 counters staffed).

TECHNOLOGY — KYC platform forces triple data entry; screen response degraded 2s to 8s.

MATERIAL/INFORMATION — app checklist contradicts branch signage; 34% arrive without documents.

ENVIRONMENT — 18 seats for a 58-minute wait; standing customers escalate faster.

MEASUREMENT — transactions-per-hour KPI with no experience measure; no branch-to-HO feedback loop.

KEY TEACHING POINTS:

1. THE HEADLINE: Head Office trained the PEOPLE bone, which carries the fewest evidenced causes. That is precisely why complaints did not move and S$ was wasted. Learners should leave able to say: 'we intervened on the wrong bone.'

2. CROSS-BONE CAUSATION: the transactions-per-hour KPI (Measurement) drives tellers to rush or avoid complex cases (Process), which pushes complex customers back into the queue (Process), which lengthens the wait (Environment), which raises complaint volume. Bones are NOT independent — this observation is the bridge into System Loops in Activity 5.

3. FISHBONE vs 5 WHYS: 5 Whys gave one deep chain; Fishbone gives six shallow ones. Depth without breadth means you fix one cause and the problem persists; breadth without depth means you fix symptoms everywhere. Mature practice uses both — and the deliberate contrast between Activity 2 and Activity 3 is the point of running them back to back.

4. GenAI is excellent at populating bones fast (breadth at near-zero cost) and poor at knowing which are real. Note how many causes it returned as [HYPOTHESIS] — those are the ones your team must go and verify. The AI generates; the human evidences.

Self-Check

Your fishbone is complete when: every bone has at least three entries; every entry is marked [EVIDENCED] or [HYPOTHESIS]; every entry is phrased as a cause ('single queue with no triage') rather than a complaint ('the queue is terrible') or a solution ('add more counters'); and you have found at least one cross-bone arrow. If any bone is empty, you have a visibility gap, not an absence of causes — say so explicitly.

Activity 4 — Pareto Analysis to Target the Vital Few

Nexa Logistics — 2,000 failed deliveries a month

Topic	Topic 1
Objective	A1, A3 · K2 — Prioritise causes by contribution and identify key implications.
Duration	40 minutes
Grouping	Teams of 3-5
Tools	Pareto Chart ed-tool, ChatGPT / Copilot / Gemini
Ed-tool	Pareto Chart Tool — https://alfredang.github.io/paretochart/

The Scenario

Nexa Logistics runs last-mile delivery for e-commerce clients across Singapore. Of 40,000 monthly deliveries, 2,000 fail on first attempt — a 5% failure rate against a contractual SLA of 2%. Every failed delivery costs S$8.50 to re-attempt: S$17,000 a month, and two clients have invoked SLA penalty clauses.

The Operations Director has a list of eleven failure reasons and a budget that will fund roughly two initiatives this quarter. His instinct is to start with 'Address incorrect' because it generates the angriest client emails.

Your team must decide where the two initiatives should go — using the data, not the noise.

The Evidence

Nexa Logistics — first-attempt delivery failures by reason, last month

Failure reason	Count	Cost impact (S$)
Recipient not at home (no delivery window given)	742	6307
Driver could not access condo / office lobby	486	4131
Address incomplete or incorrect	231	1964
Parcel not loaded (sorting error)	188	1598
Recipient refused (wrong item expected)	121	1029
Vehicle breakdown / route disruption	94	799
Weather-related suspension	61	519
Recipient contact number invalid	38	323
Parcel damaged in transit	24	204
Restricted delivery hours at destination	20	170
Other / unclassified	15	128

Discussion Questions

16. Build the Pareto chart. How many of the eleven reasons account for roughly 80% of failures? What does that ratio mean for how the Operations Director should spend his budget?
17. The Director wants to start with 'Address incorrect' because it generates the angriest emails. What does the data say, and how would you make that argument to him without dismissing his concern?
18. Calculate: if the top two reasons drop by 60%, what is the new failure rate? Does that meet the 2% SLA? Show your arithmetic.
19. The top two reasons — 'not at home' and 'lobby access' — look like customer problems. Reframe each as something Nexa controls. What changes about who owns the fix?
20. Pareto ranks by FREQUENCY. Name a situation in your own workplace where the most frequent problem is not the most important one. What would you rank by instead?

Trainer Debrief

EXPECTED ANALYSIS:
 Not at home 742 (37.1%, cum 37.1%)
 Lobby access 486 (24.3%, cum 61.4%)
 Address incorrect 231 (11.6%, cum 73.0%)
 Parcel not loaded 188 (9.4%, cum 82.4%) <- the 80% line falls here
 The remaining 7 reasons together account for under 18%.

So FOUR of eleven reasons carry 82% of the failures — and the top TWO alone carry 61%.

ARITHMETIC: a 60% reduction on the top two removes 0.6 x 1,228 = 737 failures, leaving 1,263 of 40,000 = 3.2%. That is a large win but it does NOT reach the 2% SLA (which needs failures under 800). Reaching SLA requires the top two AND a meaningful dent in reasons 3 and 4. Learners who claim the SLA is met have not done the arithmetic — make them show it.

KEY TEACHING POINTS:
1. 'Address incorrect' is only 11.6% — third place. It is loud, not large. The Director is responding to escalation volume, which correlates with client temperament, not failure frequency. The Pareto chart is how you have that conversation with evidence rather than opinion. This is the transferable political skill in this activity.
2. REFRAMING OWNERSHIP is the deeper lesson: 'recipient not at home' is only a customer problem if you never offered a delivery window. 'Lobby access denied' is only a customer problem if you never negotiated building access protocols. Both become Nexa-controllable the moment you restate them — and 61% of the problem changes owner.
3. LIMITATION: Pareto ranks by frequency, and frequency is not severity. A single safety-critical or regulatory failure can outrank 700 inconveniences. Always ask what the right ranking dimension is — count, cost, risk or customer harm. Here, cost tracks count closely

(uniform S$8.50), which is why the two rankings agree; that is not always true.
4. Pareto tells you WHERE to look, never WHY. The 742 'not at home' failures still need a 5 Whys. The tools chain: Pareto for priority, then 5 Whys for cause.

Self-Check

Your analysis is sound when: the cumulative column reaches 100%; you can name the vital few and the exact cumulative percentage at your cut-off; you have shown the arithmetic on whether the 2% SLA is reachable (it is not, on the top two alone); and you can state in one sentence what Pareto does NOT tell you. Verify the AI's percentages against the tool's — if they disagree, the tool is right and you have just learned something about trusting AI arithmetic.

Activity 5 — System Loops: Why the Problem Keeps Coming Back

Horizon Bank revisited — the fix that made it worse

Topic	Topic 1
Objective	A3 · K2 — Deduce linkages and patterns to identify key implications on organisational systems.
Duration	50 minutes
Grouping	Teams of 3-5
Tools	System Loop ed-tool, ChatGPT / Copilot / Gemini
Ed-tool	System Loop Tool — https://alfredang.github.io/systemloop/

The Scenario

Return to Horizon Bank Jurong East. Acting on the fishbone, the Regional Director did the obvious thing: he ordered two additional counters staffed at peak, redeploying staff from the back office.

Month 1: average wait fell from 58 to 34 minutes. Everyone celebrated.
Month 4: average wait is back to 61 minutes — WORSE than before. Complaints are at 68 per month. Three more tellers have resigned. Back-office processing is now two days behind, which has started generating a new complaint category of its own.

The Regional Director is baffled: "We added capacity and the problem came back bigger. What am I missing?"

What he is missing is that a branch is not a queue — it is a system of feedback loops. Your team must map them.

The Evidence

Horizon Bank — system variables and observed 4-month movements

Variable	Month 1	Month 4	Direction
Counters staffed at peak	3.2	5.1	up
Average wait time (min)	58 -> 34	61	down then up
Back-office processing backlog (days)	0.5	2.0	up
Teller errors requiring rework	6%	14%	up
Rework returning to the counter	low	high	up
Teller overtime hours / month	88	196	up
Staff resignations	2	3	up

(quarter)			
Share of counter staff under 6 months	44%	58%	up
Mean transaction handling time (min)	9	13	up
Customer complaints / month	60	68	up
Digital channel adoption	21%	20%	flat

Discussion Questions

21. Map at least two reinforcing (R) and two balancing (B) loops. Which loop did the Regional Director believe he was pulling, and which loop did he actually trigger?
22. Month 1 showed real improvement. Explain the DELAY: why did the harmful loop take three months to overwhelm the helpful one, and what does that mean for how we evaluate quick wins?
23. Follow the back-office redeployment all the way round. Trace how a decision about counters ended up increasing teller resignations.
24. Digital adoption stayed flat at 20% throughout. Which loop is NOT running that should be, and why is a dormant balancing loop as dangerous as an active reinforcing one?
25. Name two leverage points. What makes a leverage point different from simply doing more of the same thing harder?

Trainer Debrief

EXPECTED LOOPS:

B1 — CAPACITY CONTROL (balancing, what the Director intended): wait time up (+) -> management pressure (+) -> counters staffed (+) -> service capacity (+) -> wait time DOWN (-). This is the loop that produced the month-1 win. It is real, but it is not the only loop running.

R1 — REDEPLOYMENT BACKFIRE (reinforcing, the loop he triggered without seeing it): counters staffed (+) -> back-office staff removed (+) -> processing backlog (+) -> errors and unresolved cases (+) -> rework returns to the counter (+) -> counter workload (+) -> handling time (+) -> WAIT TIME (+) -> more pressure to staff counters (+). Handling time rising from 9 to 13 minutes is the fingerprint of this loop.

R2 — BURNOUT SPIRAL (reinforcing): workload (+) -> overtime (+) -> stress and burnout (+) -> resignations (+) -> share of inexperienced tellers (+) -> error rate (+) -> rework (+) -> workload (+). Tellers under 6 months rising 44% -> 58% while errors doubled 6% -> 14% is the fingerprint here.

B2 — DIGITAL MIGRATION (balancing, DORMANT): wait time (+) -> willingness to try digital (+) -> digital adoption (+) -> branch transaction volume (-) -> wait time (-). Adoption stayed flat at ~20%, so this loop never engaged. Nothing was done to activate it.

KEY TEACHING POINTS:
1. THE CENTRAL INSIGHT: he pulled B1 and unknowingly powered R1 and R2. Because R-loops compound, they eventually overwhelm a one-off capacity increase. Adding capacity to a system with an active reinforcing loop buys a temporary win and a bigger relapse.
2. DELAY IS THE TRAP: B1 acts in days (staff a counter, wait drops). R1 and R2 act over months (backlog builds, errors accumulate, people quit). The delay is precisely why the month-1 result validated a decision that was making things worse. Every learner should leave able to say: 'a quick win measured too early is indistinguishable from a mistake.' This is the direct bridge into Topic 3's leading vs lagging indicators.
3. DORMANT LOOPS: B2 was available the whole time and nobody activated it. Shifting simple transactions to digital reduces the load that feeds R1. Asking 'which balancing loop is not running?' is often more productive than adding effort to loops that are.
4. LEVERAGE vs EFFORT: more counters = effort (you must sustain it forever, and it feeds R1). Leverage = breaking a link in a reinforcing loop: triage complex cases away from the counter; eliminate the triple data entry so error rate falls; protect back-office capacity so rework never returns to the counter; activate B2 with in-branch digital ambassadors. Leverage changes the structure; effort fights the structure.
5. This is why Topic 1 ends here. 5 Whys gave depth, Fishbone gave breadth, Pareto gave priority — and System Loops explains why well-intentioned fixes rebound. Learners now have the full diagnostic set before touching solutions in Topic 2.

Self-Check

Your loop map is valid when: every loop CLOSES (the last variable feeds back to the first); every link carries a (+) or (-); each loop is correctly typed as R or B by counting negative links; you can explain the month-1-to-month-4 reversal by naming which loop overwhelmed which; and your leverage points break a LINK in a reinforcing loop rather than adding more effort to a balancing one. If your answer is 'hire more people', you have described effort, not leverage.

Activity 6 — Divergent Ideation with GenAI (Breadth Before Judgement)

Meridian Health — solving the retention root cause

Topic	Topic 2
Objective	A5 · K4 — Generate a wide solution set using appropriate problem-solving techniques.
Duration	45 minutes
Grouping	Teams of 3-5
Tools	ChatGPT / Copilot / Gemini, Pinboard
Ed-tool	Pinboard — https://alfredang.github.io/pinboard/

The Scenario

Your 5 Whys in Activity 2 established the root cause at Meridian Health: a workforce retention and onboarding failure that leaves 40% of the 07:00 roster staffed by untrained relief phlebotomists at 14 minutes per draw instead of 6.

The Clinical Director has accepted the analysis and killed the S$1.2m seventh-centre proposal. She now wants options — and she has been explicit: "Don't bring me three sensible ideas. Bring me twenty, including the ones that sound ridiculous, and then tell me which four are real."

Budget guidance: up to S$250,000 over 12 months. Constraint: MOH clinical protocols on phlebotomy competency cannot be relaxed under any circumstances.

The Evidence

Meridian Health — solution constraints and context

Constraint	Detail
Budget	Up to S$250,000 over 12 months
Clinical	MOH phlebotomy competency standards are non-negotiable
Contract	07:00-09:30 fasting draw window is contractual with corporate clients
Courier	Lab courier at 09:45 is fixed by the external lab partner
Workforce	Permanent phlebotomist turnover 44%/yr; relief pool is agency-supplied
Training	Current onboarding is 2 days shadowing, no structured sign-off
Market	Phlebotomist salaries are within 5% of the market median
Shift	Early shift (06:30 start) attracts no differential pay

Discussion Questions

26. Your team generated ideas manually first, then with GenAI. Compare the two lists — how many did the AI produce that your team would never have said out loud, and why not?
27. The AI was asked for two ideas borrowed from another industry. Which cross-industry idea is most transferable to Meridian, and what would have to be true for it to work?
28. Apply SCAMPER's 'Reverse' and 'Eliminate' to the fixed 07:00-09:30 window. Which supposedly fixed constraint turns out to be an assumption rather than a real constraint?
29. Identify one idea on your list that is genuinely ruled out by the MOH clinical constraint. How do you tell a hard constraint from a habit?
30. Cognitive fixedness is the tendency to reach for the familiar solution. Find one idea on your list that only appeared because you deliberately broke fixedness — what triggered it?

Trainer Debrief

EXPECTED RANGE — a strong list spans all five groups and includes at least these shapes:

PEOPLE & RETENTION — early-shift differential pay; career ladder to senior phlebotomist; convert best relief staff to permanent; retention bonus at 12 and 24 months; exit-interview programme to find the real reason people leave.

TRAINING & ONBOARDING — structured 10-day onboarding with competency sign-off; simulation arm practice before live draws; buddy system pairing relief with permanent staff; a micro-credential the agency pool can earn before their first shift.

PROCESS & SCHEDULING — stagger appointment slots so relief staff take simpler draws; roster all 5 permanent staff at 07:00 and taper later; pre-screen difficult-draw patients to trained staff; negotiate a second courier pickup.

TECHNOLOGY — vein-finder devices to cut the 18% repeat rate; digital pre-registration; automated fasting-compliance reminders the night before; queue app.

CLIENT EXPERIENCE — corporate on-site draws at the client's office; extend the window by agreement; tiered booking so the largest client gets trained staff.

KEY TEACHING POINTS:

1. DIVERGENCE MUST PRECEDE JUDGEMENT. Teams that evaluate while generating typically stop at 6-8 ideas, all conventional. The instruction to include 'ridiculous' ideas exists because the ridiculous ones are where fixedness breaks — and the merely unusual ones next to them are often viable. Count your team's manual list against the AI list; the gap IS the lesson.
2. GENAI IS A DIVERGENCE ENGINE. This is its strongest contribution to problem solving: breadth at near-zero cost, free of the political self-censorship that stops a junior employee proposing 'pay the early shift more' in front of a Finance Director. The AI has no career risk. Note that it also has no judgement — which is Activity 7's job.
3. THE SCAMPER REVEAL: the biggest unlock is that the 07:00-09:30 window is fixed only because draws happen AT THE CENTRE. 'Reverse' (bring the draw to the client's office) and 'Eliminate' (remove the centre visit for corporate clients entirely) both dissolve the

constraint rather than working within it. Meanwhile the MOH competency standard is a genuine hard constraint — it is a regulation, not a habit. Teaching learners to separate REGULATION from HABIT is the durable skill here.

4. CROSS-INDUSTRY TRANSFER: mobile phlebotomy mirrors mobile blood-donation drives; the buddy/competency-card model comes from aviation and F&B; surge differential pay comes from ride-hailing. Analogy is a recognised problem-solving strategy, and GenAI is unusually good at it because it has read across every domain.

5. Do NOT let teams rank yet. Ranking is Activity 7. Holding the line between divergence and convergence is itself the discipline being taught.

Self-Check

Your divergence is sufficient when: you have at least 20 distinct ideas spanning all five groups; at least 4 would raise an eyebrow in a management meeting; at least 2 are borrowed from another industry and you can name it; you have one idea per SCAMPER letter; and you have correctly separated at least one HARD CONSTRAINT from at least one HABIT. If every idea on your board is comfortable, you have not diverged — you have listed.

Activity 7 — Converging: Impact-Ease Matrix and Weighted Decision Matrix

Meridian Health — choosing four from twenty

Topic	Topic 2
Objective	A5 · K4, K6 — Shortlist and evaluate the most viable ideas using decision-making techniques.
Duration	50 minutes
Grouping	Teams of 3-5
Tools	ChatGPT / Copilot / Gemini, Pinboard
Ed-tool	Pinboard — https://alfredang.github.io/pinboard/

The Scenario

You now hold twenty-plus candidate solutions from Activity 6. The Clinical Director will fund approximately four, inside S$250,000 over 12 months.

Two members of your team are already advocating loudly for their favourite ideas and the discussion is circling. This is exactly the moment where teams either make an evidence-based decision or default to whoever is most senior or most persistent.

Your job is to convert the board into a defensible shortlist using two complementary tools: a fast Impact-Ease screen, then a weighted decision matrix on the survivors.

The Evidence

Meridian Health — agreed decision criteria and weightings

Criterion	Weight	Scoring guidance (1-5)
Impact on root cause (retention/training)	35%	5 = directly fixes the root cause
Speed to measurable effect	20%	5 = effect visible within 8 weeks
Cost within S$250k envelope	20%	5 = under S$25k
Clinical/regulatory risk	15%	5 = no MOH implication at all
Staff and client acceptance	10%	5 = actively welcomed by both

Discussion Questions

31. Place all your ideas on the Impact-Ease matrix. Which quadrant is most crowded, and what does a crowded Fill-In quadrant tell you about how your team was thinking?
32. Run the weighted matrix. Did the top-scoring solution match your team's gut favourite from Activity 6? If not, what did the weighting expose?
33. Change the weight on 'Speed to measurable effect' from 20% to 40% and re-rank. Does the winner change? What does that sensitivity tell you about how robust your recommendation is?

34. The AI recommended a portfolio of four. Explain why the best four INDIVIDUAL scores are not automatically the best four TOGETHER.
35. Consensus fatigue is when a team accepts a weak option just to end the discussion. Where in this activity was your team most at risk of it, and what stopped you?

Trainer Debrief

EXPECTED SHAPE OF A STRONG ANSWER:
QUICK WINS — roster all 5 permanent phlebotomists at 07:00 and taper later (near-zero cost, immediate effect); early-shift differential pay; pre-visit fasting-compliance reminders; route difficult draws to trained staff.
BIG BETS — structured 10-day onboarding with competency sign-off; convert top relief staff to permanent; corporate on-site draws.
FILL-INS — better signage, updated leaflets, minor booking-form changes.
THANKLESS — a new centre; a full HR system replacement.

A defensible portfolio of four usually pairs an immediate stabiliser with a root-cause fix, e.g.: (1) re-roster the 07:00 shift — buys relief NOW at almost no cost; (2) early-shift differential — attacks why people leave the early shift; (3) structured onboarding with competency sign-off — attacks the 14-min relief draw time and the 18% repeat rate; (4) convert top relief staff to permanent — converts the agency pool into retained capacity.

KEY TEACHING POINTS:
1. TWO TOOLS, TWO JOBS. Impact-Ease is a fast screen that clears the board of Fill-Ins and Thankless Tasks in minutes. The weighted matrix is the rigorous instrument you use on the survivors and, crucially, the one you can DEFEND to a board. Using the heavy tool on all twenty wastes the session; using only the light tool leaves you unable to justify the call.
2. THE WEIGHTS ARE THE REAL DECISION. Most teams argue about scores when the disagreement is actually about weights. Surfacing weights first — before any scoring — converts a political argument into an explicit, recorded choice. This is the single most useful facilitation move in the activity.
3. SENSITIVITY TESTING: if doubling the speed weight flips the winner, your recommendation is fragile and you must say so. If the same solution wins under several weightings, you have a robust recommendation. Boards trust the second kind. Teach learners to run the test BEFORE they present, not after they are challenged.
4. PORTFOLIO, NOT LEADERBOARD: the top four scores may all attack the same cause and leave another exposed, or two may be redundant (on-site corporate draws largely removes the need for the 07:00 re-roster for those clients). A portfolio balances quick stabilisation against durable root-cause repair. This is where human judgement outperforms the AI's ranking, and learners should see that clearly.
5. AI LIMITATION TO NAME OUT LOUD: the AI will confidently invent cost figures and timelines it has no basis for. Its ranking is only as good as the criteria you set. It is a fast, tireless, unbiased-by-politics scorer — not a decision maker. The Clinical Director will ask YOU why, not the model.

Self-Check

Your shortlist is defensible when: every surviving idea sits in a named quadrant; the weighted matrix shows criteria, weights, scores and totals; you have run at least one sensitivity test and can state whether the winner held; your four are a portfolio rather than the top four scores; the total cost fits inside S$250k; and you can name one thing the portfolio does NOT fix. If you cannot explain to the Clinical Director why idea #5 lost, the matrix has not done its job.

Activity 8 — Building the Corrective Action Plan

Meridian Health — from chosen solutions to an auditable plan

Topic	Topic 2
Objective	A4 · K9 — Develop corrective action plans for shortfalls identified.
Duration	45 minutes
Grouping	Teams of 3-5
Tools	ChatGPT / Copilot / Gemini, Pinboard
Ed-tool	Pinboard — https://alfredang.github.io/pinboard/

The Scenario

The Clinical Director has approved your four solutions. She now asks the question that separates a good analysis from a delivered result:

"Who is doing what, by when, with what money, and how will I know it worked?"

She also adds a warning from experience: "The last improvement project we ran had a beautiful slide deck and no owner. Six months later nothing had changed and everyone assumed someone else was doing it."

Your team must convert four approved solutions into a corrective action plan that would survive an internal audit — and that a colleague could execute without you in the room.

The Evidence

Corrective Action Plan — mandatory components (K9)

Component	Why it exists	Failure if missing
Root cause addressed	Ties the action to the diagnosis	Solution drifts to a symptom
Corrective action	The specific thing being done	Vague intent, no execution
Owner (named person)	Single point of accountability	Everyone assumes someone else
Timeline / milestone	Makes progress checkable	Slips indefinitely, unnoticed
Resources required	Budget, people, tools committed	Stalls at first resource conflict
Success measure + baseline	Defines what 'worked' means	Cannot evaluate; endless debate
Risk and mitigation	Anticipates what breaks it	Surprised by predictable failure
Review checkpoint	Forces a decision to continue/stop	Runs on past the point of failure

Discussion Questions

36. Complete all eight components for each of your four actions. Which component did your team find hardest to fill honestly, and what does that reveal?
37. Every owner must be a named role, not a department. Why does 'Operations' fail as an owner where 'Centre Operations Manager' succeeds?

38. Check your success measures against the baseline in your original problem statement. Can each one actually be measured with data Meridian already collects? Which cannot?
39. The AI flagged dependencies between actions. Which two of your actions must be sequenced, and what breaks if you run them in parallel?
40. Identify the biggest risk across your whole plan. Is your mitigation a real action with an owner, or a hope?

Trainer Debrief

WORKED EXAMPLE — one complete row to standard:

Root cause addressed: Untrained relief staff on the 07:00 roster (14-min vs 6-min draws).

Corrective action: Implement a structured 10-day phlebotomy onboarding with formal competency sign-off before any unsupervised draw.

Owner: Clinical Training Lead.

Timeline: Design by week 3; pilot at two centres weeks 4-8; full rollout by week 14.

Resources: S$65,000 (trainer time, simulation arms, backfill cover); 0.4 FTE for 14 weeks.

Success measure: Mean relief-staff draw time from 14 min (baseline) to under 8 min by week 16; first-attempt success from 82% to 92%.

Key risk: Backfill cover unavailable during pilot, so centres refuse to release staff. Mitigation: Clinical Director pre-commits agency budget for pilot weeks; escalation path named.

Review checkpoint: Week 8 — decision to roll out, extend the pilot, or stop.

KEY TEACHING POINTS:

1. THE OWNER IS THE PLAN'S SPINE. 'Operations' cannot be phoned, cannot be held to a date and cannot be asked why it slipped. A named role can. The Clinical Director's warning in the scenario is the most common real-world failure of improvement projects, and it is entirely preventable with one discipline: one named owner per action, no exceptions, no co-owners.
2. THE MEASURE MUST TIE TO THE BASELINE FROM TOPIC 1. This is where the course closes its own loop: the problem statement's baseline in Activity 1 is what makes evaluation possible in Topic 3. A team that skipped the baseline now literally cannot write this field — let them discover that themselves; it lands harder than being told.
3. MEASURABILITY CHECK: 'improved staff morale' is not measurable with data Meridian holds. 'Voluntary resignations per quarter' and 'early-shift roster fill rate' are. Force the substitution now, not at review time.
4. SEQUENCING: converting relief staff to permanent DEPENDS on the onboarding programme existing — convert first and you have permanent staff trained to the old inadequate standard, locking in the defect. Dependencies are where good plans quietly fail.
5. OWNER OVERLOAD: if the Centre Operations Manager owns three of four actions, the plan is a single point of failure regardless of how good each row looks. GenAI is genuinely useful at spotting this pattern mechanically — a good example of AI catching what a team misses because they are too close to it.
6. THE REVIEW CHECKPOINT MUST CARRY A DECISION. 'Review progress' is theatre. 'Decide to roll out, extend or stop' is governance.

Self-Check

Your CAP is audit-ready when: all eight components are complete for all four actions; every owner is a named role and no role owns more than two actions; every success measure has a metric, a baseline number and a target number; every measure is obtainable from data the organisation already has (or has an explicit task to start collecting it); dependencies are sequenced; every mitigation is an action with an owner; and every checkpoint names a decision. The real test: hand the plan to another team and ask them to execute row one without asking you a single question.

Activity 9 — Implementation Planning and Stakeholder Resistance

ShopFront SG — the rollout that has to survive contact with people

Topic	Topic 3
Objective	A6 · K3, K8 — Draw up implementation plans and address factors affecting effectiveness.
Duration	50 minutes
Grouping	Teams of 3-5
Tools	ChatGPT / Copilot / Gemini, Pinboard
Ed-tool	Pinboard — https://alfredang.github.io/pinboard/

The Scenario

Return to ShopFront SG from Activity 1. The diagnosis is complete and the solution set is approved: a one-page guest checkout, a CDN and caching layer to bring app load under 2 seconds, a supplier QC checklist to cut the return rate, and cart-recovery notifications.

The Chief Operating Officer has seen improvement projects die before. She says:

"The plan is fine. Now tell me who is going to fight it, why, and what you are going to do about it. Last year we bought a S$400,000 CRM that nobody used. It worked perfectly. That was not the problem."

Your team must build the implementation plan AND the stakeholder strategy that makes it stick. Assume nothing about goodwill.

The Evidence

ShopFront SG — stakeholder map and known positions

Stakeholder	Interest	Likely position
IT Development team	Already at capacity on a POS migration	Resist — sees added workload
Head of Marketing	Sponsored the S$180k discount campaign	Resist — solution implies her spend was wrong
Store Operations	In-store sales flat, feels unaffected	Indifferent — 'not my problem'
Suppliers (12 vendors)	Face new QC standards and possible delisting	Resist — cost and scrutiny
Customer Service	Handling 410 tickets/week, overloaded	Support — expects relief
Finance	Approved budget, wants ROI evidence	Conditional — needs measurement
COO (sponsor)	Accountable for the turnaround	Champion
Frontline retail staff	Fear digital shift reduces store headcount	Anxious — job security

Discussion Questions

41. Build the implementation plan for all four solutions. Which two MUST be sequenced rather than run in parallel, and what breaks if you ignore that?
42. The Head of Marketing resists because the diagnosis implies her S$180k campaign was misdirected. This is a face problem, not a logic problem. How do you secure her cooperation without requiring her to admit error?
43. Frontline retail staff fear the digital push costs them jobs. Is that fear irrational? What would you actually commit to — and what happens to your credibility if you promise something you cannot guarantee?
44. Twelve suppliers face new QC standards. What is the difference between imposing the checklist and getting them to adopt it, and which one survives after month three?
45. Identify the single stakeholder most likely to sink this rollout. Note: it may not be the loudest one. Justify your choice.

Trainer Debrief

EXPECTED IMPLEMENTATION SEQUENCING:
Guest checkout (weeks 1-4, quick win — builds momentum and credibility with Finance).
CDN/caching (weeks 2-10, big bet — MUST precede or accompany cart-recovery notifications; driving traffic back to a 4.8-second app converts a recovery campaign into a second bad experience and burns the customer twice).
Supplier QC (weeks 3-16, longest lead time — external parties, contractual change).
Cart recovery (weeks 10-14, AFTER load time is fixed).

STAKEHOLDER STRATEGY — the key moves:
MARKETING is the classic trap. The objection is not logical, it is reputational. The move is to reframe rather than relitigate: the campaign proved demand exists; conversion is where it leaks. Give her ownership of the cart-recovery workstream so she wins visibly from the fix. Delivered by the COO, peer to peer — never by the analyst who produced the diagnosis.
IT DEVELOPMENT resists capacity, not the idea. The move is real relief, not encouragement: re-sequence the POS migration, or fund contract capacity. Delivered by the COO because only the sponsor can re-prioritise.
SUPPLIERS respond to incentives, not instructions. Co-design the checklist with the three best-performing vendors, then tie compliance to order volume. An imposed checklist is complied with on paper for two months; a co-designed one with commercial consequence sticks.
FRONTLINE STAFF have a rational fear and deserve an honest answer. Commit to what you can actually control — no headcount reduction tied to this programme, and involvement in click-and-collect roles. Never promise what you cannot guarantee; one broken promise here costs you every subsequent change programme.
STORE OPERATIONS indifference is the quiet risk — click-and-collect needs them and they have no reason to care yet. Give them a stake before you need them.

MOST DANGEROUS STAKEHOLDER — accept a well-argued case for either, but push teams past the obvious: IT Development can hard-block delivery (nothing ships without them), while Marketing can soft-block through political attrition, which is harder to see and harder to escalate. Many teams name the loudest resister; the more sophisticated answer notes that INDIFFERENCE (Store Operations) often kills rollouts more quietly than opposition, because nobody escalates an absence of enthusiasm.

KEY TEACHING POINTS:
1. THE COO'S LINE IS THE WHOLE LESSON: the CRM 'worked perfectly. That was not the problem.' Solutions fail on ADOPTION far more often than on design. K8 — factors affecting the effectiveness of an implementation plan — is mostly about people, not technology.
2. MESSENGER MATTERS AS MUCH AS MESSAGE. The same words land differently from the COO, a peer, or the analyst. Teams consistently under-think this; make them assign a messenger to every message and justify it.
3. RESISTANCE IS INFORMATION. IT's objection reveals a genuine capacity conflict your plan had not costed. Suppliers' objection reveals the checklist needs commercial teeth. Treating resistance as data improves the plan; treating it as obstruction guarantees failure.
4. GenAI is unusually good here because it has no stake in the office politics — it will state plainly that Marketing's objection is about face, which a junior analyst may not dare write down. It is weak on the specific human relationships and history, which only the team knows. Divide the labour accordingly.

Self-Check

Your implementation plan is realistic when: every solution has an action, method, owner and timeline; dependencies are sequenced (cart recovery AFTER load time is fixed); every stakeholder has an objection written in their own voice, a specific move, and a named messenger; you have identified the highest-risk blocker with justification; and your communication plan runs the full 90 days rather than stopping at launch. If your plan assumes everyone cooperates because the analysis is correct, you have written a wish, not a plan.

Activity 10 — Measuring Effectiveness: Did It Actually Work?

ShopFront SG — 90 days later, and the numbers are ambiguous

Topic	Topic 3
Objective	A7 · K7 — Evaluate the effectiveness of implemented solutions and implementation plans.
Duration	45 minutes
Grouping	Teams of 3-5
Tools	ChatGPT / Copilot / Gemini, Pareto Chart ed-tool
Ed-tool	Pareto Chart Tool — https://alfredang.github.io/paretochart/

The Scenario

Ninety days have passed. The COO has the results and she is not sure what to make of them. Some numbers moved well. Some did not move at all. One moved in the wrong direction. And the quarter included the Great Singapore Sale, which everyone is now using to explain whichever result suits their argument.

The Head of Marketing says the improvement is obviously seasonal. The IT lead says the improvement is obviously the CDN. Finance wants to know whether to release the next S$200,000 tranche.

Your team must evaluate honestly: what worked, what did not, what you cannot yet tell, and what you recommend next. "It's working" is not an acceptable answer.

The Evidence

ShopFront SG — 90-day results against baseline

Metric	Baseline	Target	Day 90 actual	Verdict
App page load	4.8 s	under 2.0 s	1.7 s	Target met
Checkout completion	51%	68%	63%	Improved, short
Cart abandonment	74%	60%	64%	Improved, short
Online sales / month	S$1.9m	S$2.8m	S$2.5m	Improved, short
CSAT	70%	85%	72%	Barely moved
Complaint tickets / week	410	150	295	Improved, short
Product return rate	12%	5%	13%	WORSE
Repeat purchase (90-day)	29%	40%	31%	Barely moved
Guest checkout adoption	n/a	n/a	44% of orders	New data
Supplier QC checklist	0	12 of 12	4 of 12	Behind

adoption				
Cart recovery emails sent	n/a	n/a	18,400	Live
Cart recovery conversion	n/a	8%	3.1%	Underperforming

Discussion Questions

46. Give a verdict on each metric. Overall — did this programme succeed? Defend a single-sentence answer to the COO.
47. The product return rate got WORSE (12% to 13%) despite the QC initiative. Give the most likely explanation, and say what evidence would confirm or refute it.
48. App load smashed its target but CSAT barely moved. Explain this using leading versus lagging indicators. How long would you expect CSAT to lag?
49. The Great Singapore Sale sits inside the measurement window. How would you isolate the programme effect from the seasonal effect? What would you have done differently at day 0 to make this answerable?
50. Cart recovery converts at 3.1% against an 8% target. Is the solution wrong, or the implementation, or the target? How would you tell the difference — and what does each answer imply?
51. Make the call to the CFO: release the next S$200,000, adjust, or stop. Justify it with the evidence, and name what would change your mind.

Trainer Debrief

EXPECTED EVALUATION:

ATTRIBUTION IS THE CENTRAL SKILL. App load 4.8s -> 1.7s is cleanly attributable to the CDN and caching — technical, direct, no plausible alternative cause. Checkout completion is PARTIALLY attributable to guest checkout (44% adoption is strong supporting evidence). Online sales, however, are CONFOUNDED by the Great Singapore Sale and cannot be cleanly attributed either way. A team that credits the full S$0.6m to the programme is overclaiming; a team that dismisses it as seasonal is underclaiming. The honest answer is: not determinable with this design, and here is what would have made it determinable.

THE RETURN RATE PUZZLE — the most instructive result on the board. Most likely explanation: only 4 of 12 suppliers adopted the QC checklist, so the intervention barely happened. Simultaneously, higher conversion and GSS volume brought in more first-time and discount-driven buyers, who return at a higher rate than loyal customers. So the return rate rose because of a MIX SHIFT, not because quality fell. Confirming evidence: segment return rate by supplier (adopters vs non-adopters) and by customer type (new vs repeat, full-price vs discounted). This is the moment learners see that a metric moving the wrong way does not automatically mean the solution failed — you must decompose before you conclude.

LEADING vs LAGGING: app load, checkout completion and guest checkout adoption are LEADING — they respond within days. CSAT, repeat purchase and brand perception are LAGGING — they reflect accumulated experience across multiple purchase cycles and typically lag two to three cycles, i.e. 6-9 months for a fashion retailer. Expecting CSAT to move in 90 days was a planning error, not an execution failure. The target itself was wrong.

CART RECOVERY at 3.1% vs 8%: distinguish the three diagnoses. Wrong solution (cart recovery does not work for this audience); wrong implementation (poor timing, weak copy, no incentive, emails hitting spam); wrong target (8% was benchmarked from a different sector). The way to tell: check open and click rates. High open + low conversion = offer/landing problem. Low open = deliverability or timing problem. Industry benchmarks for fashion cart-recovery sit around 3-5%, which strongly suggests the TARGET was unrealistic. Teams that immediately conclude 'the solution failed' have skipped the diagnosis — and that is exactly the behaviour the whole course exists to correct.

RECOMMENDATION TO THE CFO — the defensible answer is CONTINUE WITH ADJUSTMENT: the technical solutions delivered and are attributable; supplier QC did not actually get implemented (4 of 12) so it has not been tested and should not be abandoned on non-adoption; cart recovery needs its target rebased and its execution diagnosed; CSAT needs a longer horizon. Redirect effort to supplier adoption — which is a CHANGE MANAGEMENT failure, traceable straight back to Activity 9's prediction that suppliers would resist. That connection is the strongest single moment in the course: the resistance they predicted is the resistance that actually materialised, and it is the reason a metric missed.

KEY TEACHING POINTS:
1. 'IT'S WORKING' IS NOT AN EVALUATION. Metric-by-metric verdicts against baseline and target, with explicit attribution and explicit uncertainty, is.
2. EVALUATION IS ONLY POSSIBLE BECAUSE ACTIVITY 1 CAPTURED A BASELINE. Say this out loud — the course has now closed its loop, and learners should feel why the six-element problem statement mattered 16 hours ago.
3. DISTINGUISH SOLUTION FAILURE FROM IMPLEMENTATION FAILURE. Supplier QC did not fail; it was not adopted. Those demand completely different responses — one means change the solution, the other means fix the rollout.
4. CONFOUNDING IS ALWAYS PRESENT IN REAL WORKPLACES. You will rarely get a clean experiment. Good practice: define the measurement window at day 0, hold a control (e.g. staged rollout by region), and segment the data. Say what you cannot determine — that honesty is what makes the rest of your evaluation credible to a CFO.
5. GenAI will produce a confident, well-structured evaluation that quietly overclaims attribution unless you explicitly instruct it to flag what cannot be determined. Note that the prompt on the slide does exactly that — and compare it against what the AI returns without that instruction. That contrast is the AI-literacy takeaway of the entire course.

Self-Check

Your evaluation is credible when: every metric has a verdict against its baseline AND target; attribution is stated explicitly and confounded results are named as confounded rather than claimed; you can explain the worsening return rate without concluding the solution failed; you distinguish solution failure from implementation failure for supplier QC; your CFO recommendation names what would change your mind; and you can trace at least one missed metric back to a stakeholder resistance you predicted in Activity 9. If your evaluation credits the programme with everything that improved and blames the season for everything that did not, you have written advocacy, not evaluation.